

# INTRO TO DATA SCIENCE LESSON 3: INTRO TO DATABASES

**WHAT IS MACHINE LEARNING?**  
**TYPES OF MACHINE LEARNING PROBLEMS**  
**DESIGNING A REGRESSION USING MATRICES**  
**USING NUMPY, SCIPY, AND PANDAS**

**QUESTIONS?**

**I. INTRO TO DATABASES**

**II. RELATIONAL DATABASES / SQL**

**LAB**

**III. SQL LAB**

**IV. INTERACTING WITH SQL FROM PANDAS**

# INTRO TO DATA SCIENCE

---

## I. INTRO TO DATABASES


Databases are a **structured** data source optimized for efficient **retrieval and storage**.


Databases are a **structured** data source optimized for efficient **retrieval** and **persistent storage**.

**structured:** we have to pre-define organization strategy

**retrieval:** the ability to read data out

**storage:** the ability to write data and save it




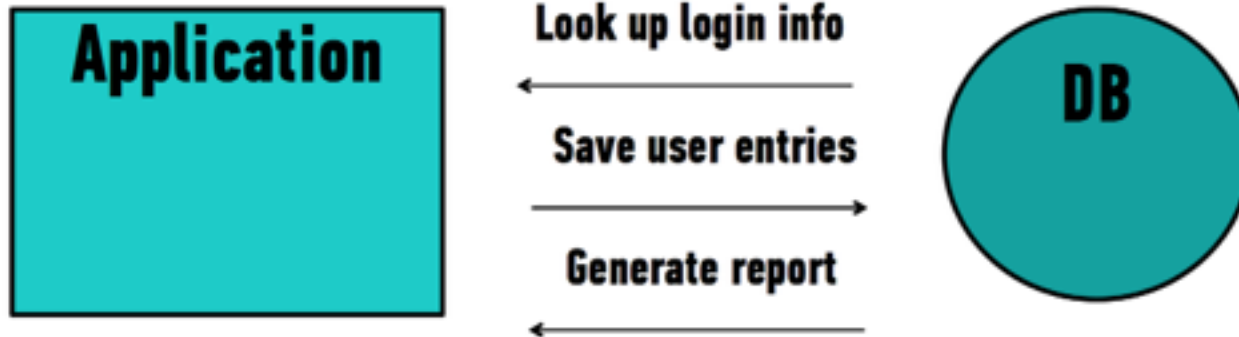


Look up login info









## INTRO TO DATA SCIENCE

---

# II. RELATIONAL DATABASES

Relational databases are traditionally organized in the following manner:

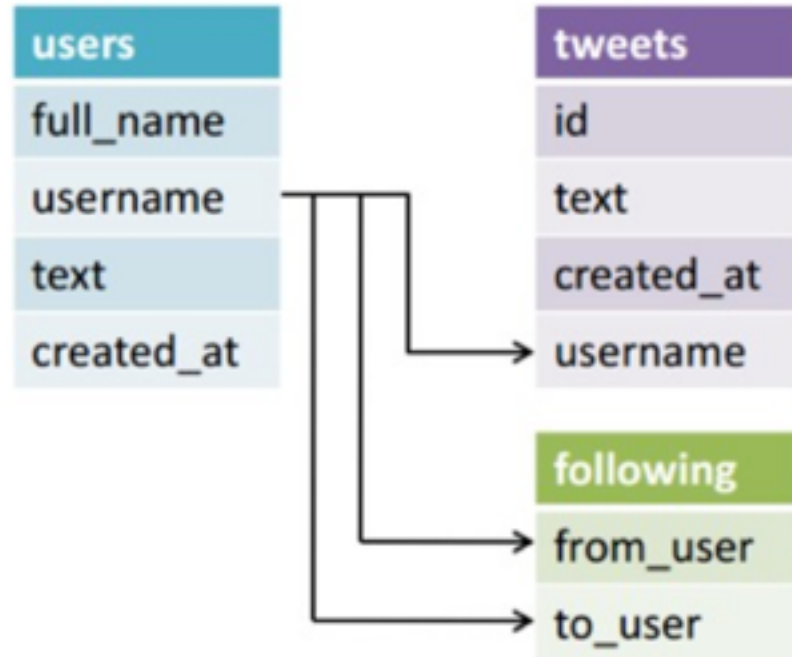
**A database has tables which represent individual entities or objects**

**Tables have predefined schema – rules that tell it what the data will look like**

Each table should have a **primary key** column – a unique identifier for that row

Each table should have a **primary key** column – a unique identifier for that row

Additionally each table can have a **foreign key** column – an id that links this to another table.



We could have had a table structure as follows:  
Why is this different?

tweets
id
text
created_at
username
full_name
username
text
created_at



We could have had a table structure as follows:

Why is this different?

We would repeat the user information in each row.

This is called  
**denormalization.**

tweets
id
text
created_at
username
full_name
username
text
created_at


**Normalized Data:** Many tables to reduce redundant or repeated data in a table

**Denormalized Data:** Wide data with fields often repeated but removes the need to join together multiple tables

**Normalized Data:** Many tables to reduce redundant or repeated data in a table

**Denormalized Data:** Wide data with fields often repeated but removes the need to join together multiple tables

**This is a trade off of speed vs storage.**

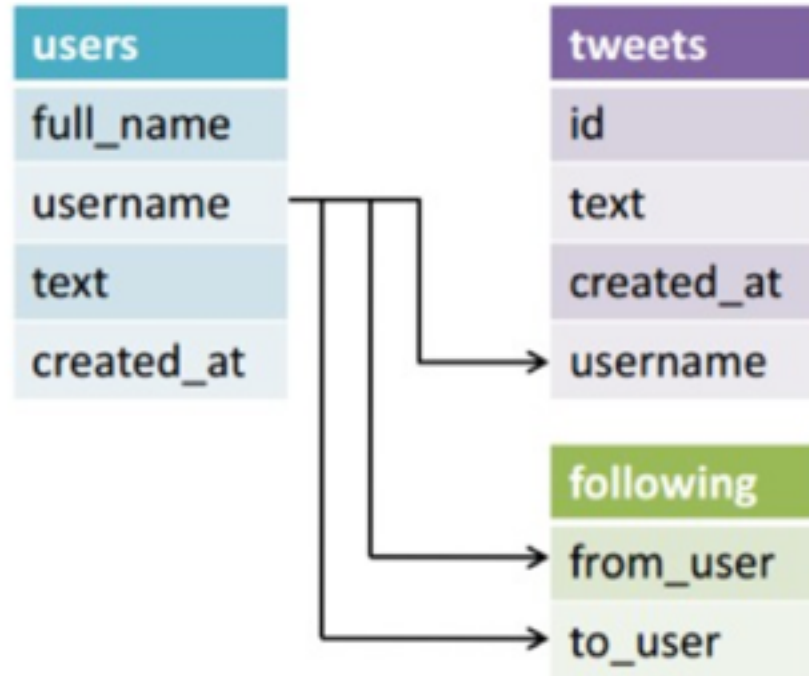


Q: How do we commonly evaluate databases?

**read-speed vs write-speed**  
**space considerations**  
**(and many, many other criteria)**

Q: Why are normalized tables (possibly) slower to read?

**A: We'll have to get data from multiple tables to answer some questions**



Q: Why are denormalized tables (possibly) slower to write?

**A: We'll have to write more information on each write**



tweets
id
text
created_at
username
full_name
username
text
created_at

SQL is a query language to load, retrieve, and update data in relational databases

Most commonly known SQL-like Databases include:

Oracle

MySQL

PostgreSQL

**SELECT:** Allows you to retrieve information from a table

Syntax:

```
SELECT col1, col2  
FROM table WHERE [some condition]
```

Example:

```
SELECT poll_title, poll_date FROM polls WHERE  
romney_pct > obama_pct
```

**GROUP BY:** Allows you to aggregate information.

Syntax:

```
SELECT col1, AVG(col2)  
FROM table GROUP BY col1
```

Example:

```
SELECT poll_date, AVG(obama_pct)  
FROM polls GROUP BY poll_date
```

**GROUP BY:** Allows you to aggregate information.

Syntax:

```
SELECT col1, AVG(col2)  
FROM table GROUP BY col1
```

Example:

```
SELECT poll_date, AVG(obama_pct)  
FROM polls GROUP BY poll_date
```

**GROUP BY:** Allows you to aggregate information.

Syntax:

```
SELECT col1, AVG(col2)
FROM table GROUP BY col1
```

There are usually a few common built-in operations:

**SUM, AVG, MIN, MAX, COUNT**

**JOIN:** Allows you to combine multiple tables

Syntax:

```
SELECT t1.c1, t1.c2, t2.c2  
FROM t1 JOIN t2 ON t1.c1 = t2.c2
```

**JOIN:** Allows you to combine multiple tables

Syntax:

```
SELECT t1.c1, t1.c2, t2.c2  
FROM t1 JOIN t2 ON t1.c1 = t2.c2
```




```
id name      id name
-- ----      -- ----
1  Pirate     1  Rutabaga
2  Monkey     2  Pirate
3  Ninja      3  Darth Vader
4  Spaghetti  4  Ninja
```

Let's join these tables by the name field in a few different ways and see if we can get a conceptual match to those nifty Venn diagrams.

```
SELECT * FROM TableA
INNER JOIN TableB
ON TableA.name = TableB.name
```

```
id name      id name
-- ----      -- ----
1  Pirate     2  Pirate
3  Ninja      4  Ninja
```

Inner join produces only the set of records that match in both Table A and Table B.



**INSERT:** Allows you to **add** data to tables

Syntax:

```
INSERT INTO table1 (col1, col2)  
VALUES (...)
```

```
INSERT INTO classroom (first_name, last_name)  
VALUES ('John', 'Doe')
```

NoSQL databases are a new trend in databases

The name **NoSQL** refers to the lack of a relational structure between stored objects.

Most importantly they attempt to minimize the need for **JOIN** operations, or solve other data needs

Memcached	::	Livejournal
Apache HBase	::	Google BigTable
Cassandra	::	Amazon Dynamo
MongoDB	::	10Gen
Hadoop	::	Google MapReduce

---

**INTRO TO DATA SCIENCE**

---

**LAB**